

# Module 7

Taylor

2022-07-05

## Schedule Information

- ▶ HW3 Due 7/17

# Outline

Two unrelated topics:

- ▶ Hierarchical models
- ▶ Dealing with model uncertainty

## Hierarchical models

Why are we talking about **hierarchical models**?

- ▶ when you need them, you need them
- ▶ HMC and NUTS have a hard time with them unless you reparameterize

## Conditional IID (aka Exchangeability)

Most setups in Bayesian ML:

- ▶  $\theta \sim \text{prior}$
- ▶  $y_1, \dots, y_n | \theta$  independent and identically distributed

Does your data violate this assumption?

- ▶ time series
- ▶ repeated measures/longitudinal data
- ▶ **clusters/groups**

# Simpson's “Paradox” (an aside)




Figure 1: source: <https://paulvanderlaken.com/2017/09/27/simpsons-paradox-two-hr-examples-with-r-code/>

# Hierarchical Regression

Regular regression:

- ▶  $y_i \mid a, b, \sigma^2, x_i \sim \text{Normal}(a + bx_i, \sigma^2)$
- ▶ prior:  $p(a, b)p(\sigma^2)$

Hierarchical Regression:

- ▶  $y_{ij} \mid a_j, b_j, \sigma^2, x_{ij} \sim \text{Normal}(a_j + b_j x_{ij}, \sigma^2)$
- ▶ prior:  $\{\prod_j p(a_j, b_j \mid \theta)\}p(\theta)p(\sigma^2)$

TLDR:

- ▶ each group gets its own intercept and slope
- ▶ overfitting is prevented by “tying together” betas in the prior

# Hierarchical Regression in PYMC3

Demo time... hierarchical\_model\_demo.ipynb

## Model Uncertainty

Changing gears now...

"All models are wrong." What do we do about this?

- ▶ Ignore model uncertainty
- ▶ Expand/revise your working model
- ▶ select one model from several (e.g. hypothesis testing, choosing one with best OOS predictions)
- ▶ Average/blend many models together

## Choosing one from many: option 1

Fit a bunch of models and pick the one that predicts the best. This is the same as picking the model with the best *scoring rule* or *information criterion* score:

$$\mathbb{E}_{\tilde{y}}[-\log p(\tilde{y})]$$

- ▶  $y$  is your data set
- ▶  $\tilde{y}$  is future data
- ▶  $p(\tilde{y})$ : frequentists use  $p(y | \hat{\theta})$ , Bayesians use  $p(\tilde{y} | y)$
- ▶ this quantity isn't even close to being tractable

## Choosing one from many: option 1

$$\mathbb{E}_{\tilde{y}}[-\log p(\tilde{y})]$$

can be approximated by

- ▶ train/test splits (e.g.  $S^{-1} \sum_{j=1}^S -\log p(\tilde{y}_j)$ )
- ▶ cross-validation (i.e. many train/test splits, usually computationally brutal)
- ▶ AIC, WAIC, BIC, DIC etc. (use all your data—no split; goodness of fit + penalty term)

## Choosing one from many: option 2

Then there's hypothesis testing:

- ▶ null hypothesis:  $M_1$  an entire model with prior  $p(\theta_1 | M_1)$
- ▶ alternative:  $M_1$  a different model with totally different set of parameters—prior:  $p(\theta_2 | M_2)$
- ▶ also have prior beliefs on which model is the true model  $P(M_1)$  and  $P(M_2)$

$$\begin{aligned}\frac{p(M_2 | y)}{p(M_1 | y)} &= \frac{p(y | M_2)p(M_2) / p(y)}{p(y | M_1)p(M_1) / p(y)} \\ &= \underbrace{\frac{p(y | M_2)}{p(y | M_1)}}_{\text{Bayes factor}} \frac{p(M_2)}{p(M_1)}\end{aligned}$$

## Choosing one from many: option 2

$$\frac{p(M_2 | y)}{p(M_1 | y)} = \underbrace{\frac{p(y | M_2)}{p(y | M_1)}}_{\text{Bayes factor}} \frac{p(M_2)}{p(M_1)}$$

model evidence:

$$p(y | M_2) = \int p(y | \theta_2, M_2) p(\theta_2 | M_2) d\theta_2$$

BFs are tricky:

- ▶ evidences are sensitive and can be difficult to calculate/approximate
- ▶ BDA3 text dismisses them and favors “continuous model expansion”

## Averaging Many Models Together

**Bayesian Model Averaging** is an overarching principle that follows from simple definitions:

$$p(\tilde{y} \mid y) = \sum_i p(\tilde{y} \mid M_i, y) p(M_i \mid y)$$

- ▶ weights:  $p(M_i \mid y)$
- ▶ predictions:  $p(\tilde{y} \mid M_i, y)$
- ▶ a good review: <https://www.jstor.org/stable/2676803>

# Averaging Many Models Together

It's difficult:

$$p(\tilde{y} \mid y) = \sum_i p(\tilde{y} \mid M_i, y) p(M_i \mid y)$$

- ▶  $p(M_i \mid y) \propto p(M_i) \int p(y \mid \theta_i, M_i) p(\theta_i \mid M_i) d\theta_i p(M_i)$
- ▶  $p(\tilde{y} \mid M_i, y) = \int p(\tilde{y} \mid \theta_i, M_i) p(\theta_i \mid M_i, y) d\theta_i$
- ▶ list out all possible models, compute posteriors for each one, then posterior predictive distributions, then prior predictive distributions...
- ▶ alternative decompositions might be more accessible via fancy MCMC sampling targeting model  $\times$  parameter space